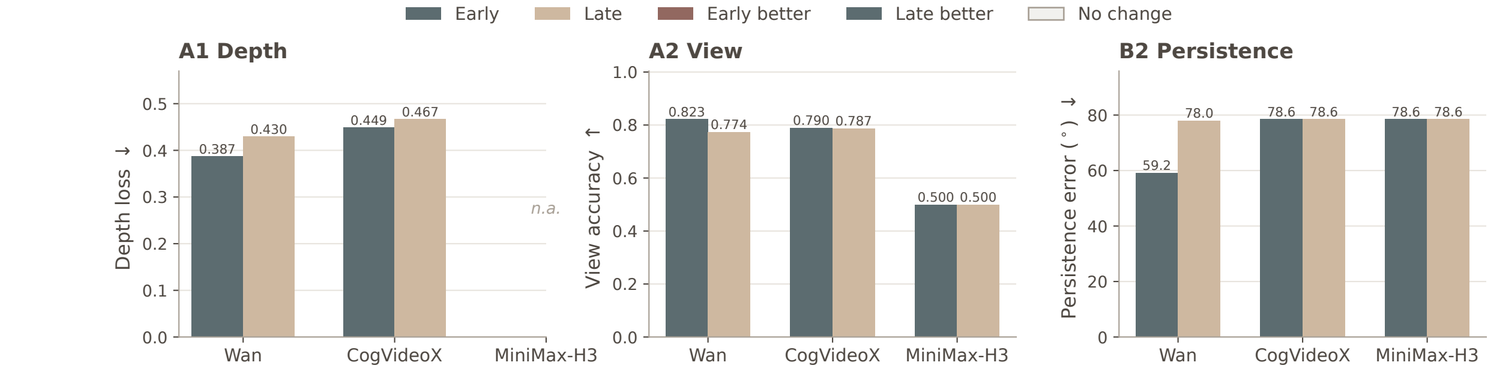

Layer–prefix pairs (E = early better, L = late better)

	L15 p8	L15 p12	L15 p16	L15 p24	L20 p8	L20 p12	L20 p16	L20 p24	L25 p8	L25 p12	L25 p16	L25 p24
Wan Depth	E	E	E	E	E	E	E	E	E	E	E	E
Wan View	E	E	E	E	E	E	E	E	E	E	E	E
Wan Persist.	E	E	E	E	E	E	E	E	E	E	E	E
CogVideoX Depth	E	E	E	L	E	E	E	E	E	E	L	L
CogVideoX View	E	E	E	E	E	L	E	E	E	E	L	E
CogVideoX Persist.	=	=	=	=	=	=	=	=	=	=	=	=
MiniMax-H3 Depth	L	L	L	L	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.
MiniMax-H3 View	E	E	E	E	=	=	=	=	=	=	=	=
MiniMax-H3 Persist.	=	=	=	=	=	=	=	=	=	=	=	=

Official held-out evaluation (eval_split=val). Each pair of bars is the same model at the early vs late scheduler index (layer 20, mean of prefixes 8/12/16/24). Lower is better for depth and persistence; higher is better for view. MiniMax-H3 depth at layers 20/25 is excluded: official loss is exactly 0. Bottom grid: E/L/= for each layer–prefix pair. CogVideoX and MiniMax-H3 persistence do not change.